



SCHOOL OF MANAGEMENT

Gautam Buddha University

Gautam Buddha Nagar – 201312, Uttar Pradesh, India

Capstone Project Report

Project Title:

AI-Based Customer Sentiment Analysis for Predicting Product Success Using -World Social Media and E-Commerce Data

Submitted in partial fulfilment of the requirements for the award of the degree of

MBA (Business Analytics and Data Science)

Batch: 2024 – 2026

By

Abhijeet Srivastava

Roll No.: 246 PBA 021

Under the Supervision of

Mr. Alok Sharma

Department of Business Analytics

May 2026

DECLARATION BY STUDENT

I, Abhijeet Srivastava, Roll No. 246 PBA 021, hereby declare that the project titled “*AI-Based Customer Sentiment Analysis for Predicting Product Success Using -World Social Media and E-Commerce Data*” is an original work carried out by me under the supervision of Mr. Alok Sharma, School of Management, Gautam Buddha University. This project has not been submitted, either in whole or in part, for any degree or diploma at this or any other institution.

All datasets used in this project are publicly available, -world Kaggle datasets — the Amazon India Product Sales Dataset, the CrowdFlower Apple-Google Tweet Sentiment Dataset, the Consumer Electronics Product Reviews Dataset, and the Customer Mobile Phone Tweet Reviews Dataset — comprising a combined corpus of 36,401 authentic records. No synthetic or mock data has been used. Any material derived from published sources has been duly acknowledged.

I am solely responsible for any errors or omissions in this work.

Date: _____

Place: Greater Noida, Uttar Pradesh

Abhijeet Srivastava

MBA (Business Analytics and Data Science) – Semester IV

Roll No.: 246 PBA 021 | School of Management, Gautam Buddha University

CERTIFICATE FROM THE SUPERVISOR

This is to certify that the project work titled “*AI-Based Customer Sentiment Analysis for Predicting Product Success Using -World Social Media and E-Commerce Data*” has been carried out under my supervision and guidance by Abhijeet Srivastava, Roll No. 246 PBA 021, a student of MBA (Business Analytics and Data Science), IV Semester, School of Management, Gautam Buddha University.

The work reported in this project is genuine and original, based on publicly available Kaggle datasets, and has not been submitted for any other degree or diploma in this or any other university or institution. The student has completed all requirements of the project satisfactorily and has fulfilled the minimum number of mentor meetings as prescribed.

Name of Supervisor: Mr. Alok Sharma

Designation: _____

Department / School: _____

Signature: _____

Date: _____

ACKNOWLEDGEMENT

I'd like to take a moment to thank the people who made this project possible.

First and foremost, I owe a lot to my supervisor, Mr. Alok Sharma, whose guidance, patience, and feedback shaped this project in ways I didn't always anticipate but always appreciated. His constructive feedback and direction were instrumental in shaping this project.

I also want to thank the Dean and faculty at the School of Management, Gautam Buddha University, for creating the kind of environment where this type of applied research is genuinely encouraged.

I am grateful to the open data community and dataset contributors whose publicly available, -world data made this research possible. Specifically, I acknowledge the contributors of: (1) the Amazon India Product Sales Dataset; (2) the CrowdFlower Apple-Google Tweet Sentiment Dataset (crowd-annotated and publicly hosted); (3) the Consumer Electronics Product Reviews Dataset; and (4) the Customer Mobile Phone Tweet Reviews Dataset. These datasets, totalling over 36,000 records, form the empirical foundation of this study.

I also acknowledge the developers of open-source tools including Python, Scikit-learn, Pandas, NumPy, Matplotlib, and Seaborn. Finally, I wish to acknowledge my family and fellow students for their continued motivation and moral support.

Abhijeet Srivastava

MBA (Business Analytics and Data Science)

School of Management, Gautam Buddha University | April 2026

TABLE OF CONTENTS

Declaration by Student	ii
Certificate from Supervisor	iii
Acknowledgement	iv
Table of Contents	v
List of Figures	vi
List of Tables	vii
List of Abbreviations	viii
Abstract	ix
Chapter 1: Introduction	1
1.1 Background of the Study	1
1.2 Introduction to the Domain	3
1.3 Scope of the Study	4
1.4 Significance of the Study	5
1.5 Type of Project	6
Chapter 2: Literature Review	7
2.1 Review of Research Papers	7
2.2 Analytical Approaches in Similar Studies	11
2.3 Conceptual Framework	12
2.4 Identified Research Gap	13
Chapter 3: Problem Statement and Objectives	14
3.1 Problem Statement	14
3.2 Research Objectives	15
3.3 Hypotheses	17
Chapter 4: Research Methodology	18
4.1 Business Understanding (CRISP-DM)	18
4.2 Data Sources	19
4.3 Data Collection Methodology	21
4.4 Tools and Platforms	22
4.5 Analytical Pipeline	23
Chapter 5: Data Preparation and Processing	24
5.1 Dataset Description and Structure	24
5.2 Class Distribution	26
5.3 Data Cleaning	27
5.4 Outlier Detection and Treatment	28
5.5 Text Preprocessing Pipeline	29
5.6 Feature Engineering	31
Chapter 6: Data Analysis, Modelling and Results	32
6.1 Exploratory Data Analysis (EDA)	32
6.2 Sentiment Analysis Using NLP (VADER)	38
6.3 Machine Learning Modelling	41
6.4 Model Evaluation and Results	43
Chapter 7: Interpretation and Business Application	49
7.1 Key Findings	49
7.2 Business Insights	51
7.3 Strategy Recommendations	53
Chapter 8: Process Flow and System Design	55
8.1 End-to-End Workflow	55

8.2 Model Pipeline Architecture	57
8.3 Model Deployment Concept	58
Chapter 9: Conclusion and Recommendations	59
9.1 Summary of Key Findings	59
9.2 Recommendations	60
9.3 Limitations	61
9.4 Future Scope	62
References	63
Appendix A: Python Code	67
Appendix B: Dataset Samples	78

LIST OF FIGURES

Figure 6.1: Overall Sentiment Distribution — Combined Dataset	32
Figure 6.2: Sentiment Distribution by Product Category	33
Figure 6.3: Amazon India Star Rating Distribution (Pie Chart)	34
Figure 6.4: Monthly Positive Sentiment Score Trend (Jan 2023 – Nov 2024)	35
Figure 6.5: Aspect-Level Sentiment Analysis — Consumer Electronics	37
Figure 6.6: Distribution of Sentiment / VADER Compound Scores	38
Figure 6.7: Normalised Confusion Matrix — Logistic Regression (Test Set)	44
Figure 6.8: ROC Curves — Logistic Regression (One-vs-Rest)	45
Figure 6.9: Top 15 Feature Importances — Positive and Negative Classes	47
Figure 6.10: ML Model Performance Comparison	48
Figure 6.11: Positive Sentiment Ratio vs Best Seller Rank Improvement	49
Figure 8.1: CRISP-DM End-to-End Analytical Pipeline	55
Figure 8.2: Model Deployment Architecture — Proposed Production Pipeline	58

LIST OF TABLES

Table 2.1: Comparative Overview of Methodologies in Sentiment Analysis Literature	11
Table 4.1: Tools and Libraries Used in the Analytical Pipeline	22
Table 5.1: Dataset Summary — Field Descriptions and Sources	24
Table 5.2: Class Distribution Across Datasets	26
Table 5.3: Amazon Review Dataset Field Descriptions	25
Table 5.4: Twitter / Tweet Dataset Field Descriptions	25
Table 6.1: VADER Sentiment Performance Metrics	38
Table 6.2: Aspect-Level Sentiment Scores Across Product Categories	40
Table 6.3: ML Model Performance Comparison on Test Set (n=1,491)	43
Table 6.4: Logistic Regression Per-Class Performance Detail	46
Table 6.5: Cross-Source Sentiment Comparison — Amazon vs. Tweets	48
Table 7.1: Business Application Framework — Findings to Actions	51
Table 8.1: End-to-End Process Flow Summary	56
Table B.1: Amazon India Dataset — Sample Records	78
Table B.2: CrowdFlower Tweet Dataset — Sample Records	79
Table B.3: Consumer Electronics Reviews — Sample Records	80

LIST OF SYMBOLS AND ABBREVIATIONS

Abbreviation	Full Form
AI	Artificial Intelligence
ML	Machine Learning
NLP	Natural Language Processing
EDA	Exploratory Data Analysis
API	Application Programming Interface
VADER	Valence Aware Dictionary and sEntiment Reasoner
TF-IDF	Term Frequency – Inverse Document Frequency
BoW	Bag of Words
LR	Logistic Regression
RF	Random Forest
NB	Naive Bayes
ROC	Receiver Operating Characteristic
AUC	Area Under the Curve
NLTK	Natural Language Toolkit
CNN	Convolutional Neural Network
LSTM	Long Short-Term Memory
MBA	Master of Business Administration
GBU	Gautam Buddha University
SoM	School of Management
CRM	Customer Relationship Management
NPD	New Product Development
VOC	Voice of Customer
KPI	Key Performance Indicator
CSV	Comma Separated Values
PSR	Positive Sentiment Ratio
BSR	Amazon Best Seller Rank
ABSA	Aspect-Based Sentiment Analysis
CRISP-DM	Cross-Industry Standard Process for Data Mining
SMOTE	Synthetic Minority Over-Sampling Technique
UGC	User-Generated Content

Table: List of Abbreviations Used in This Report

ABSTRACT

Problem Definition

"Nowadays customers share their opinions on social media and e-commerce platforms every day. These opinions can provide valuable insights for businesses. This project explores whether such data can help predict product success before traditional business metrics reflect the outcome."

Traditional product success metrics — quarterly sales reports, periodic market surveys, and post-hoc customer satisfaction studies — suffer from significant information latency: they reveal what has already happened rather than what is likely to happen next. This project tries to answer a practical question: can - world social media data give companies an early heads-up about whether a product is going to succeed or fail, before the financial numbers confirm it?

Data Source and Datasets Used

This study utilises four publicly available, Kaggle datasets comprising a combined corpus of **36,401 records**: (1) Amazon India Product Sales Dataset — 1,464 electronics product listings with authentic user ratings and review content from Amazon.in; (2) CrowdFlower Apple-Google Tweet Sentiment Dataset — 9,093 annotated tweets covering Apple and Google product sentiment; (3) Consumer Electronics Product Reviews Dataset — 1,000 structured reviews across five product categories; and (4) Customer Mobile Phone Tweet Reviews Dataset — 25,000 mobile technology tweets. All datasets are authentic, organically collected or crowd-annotated — not synthetically generated. The Amazon India dataset reports an average product rating of **4.10/5** with 75.8% of products rated Positive.

Analytical Approach and Key Findings

The pipeline brought together TF-IDF vectorisation (5,000 features, including bigrams) with three supervised ML classifiers. Three models were tested on the combined corpus (n=9,937 records) — Logistic Regression as the primary, plus Random Forest and Naïve Bayes for comparison. The Logistic Regression classifier achieved an accuracy of **73.5%** on the imbalanced test set (n=1,491), with a weighted ROC-AUC of **0.829**, significantly outperforming the majority-class baseline of 56.8%. Category-level results show Mobile Accessories leading at 86% Positive while Headphones/Earbuds lag behind at just 52%. Wearables and Laptops achieve the highest satisfaction in the Consumer Electronics Reviews dataset (62% and 61% positive respectively). The positive sentiment ratio (PSR) is shown to precede Amazon Best Seller Rank improvements by approximately 2–3 weeks, establishing its utility as a business leading indicator.

Business Application

These findings are not just interesting academically — they point to specific actions that product managers, marketing teams, and supply chain planners can take. The pipeline itself is ready to be deployed as a -time service — monitoring daily PSR by category and triggering alerts whenever key thresholds are crossed. Combining Amazon reviews with tweet sentiment also proved worthwhile — the two channels reveal meaningfully different perception patterns, and together they give a more complete picture than either would alone.

Keywords: *Sentiment Analysis, NLP, TF-IDF, Logistic Regression, Kaggle Data, Amazon India, Consumer Electronics, Product Success Prediction, Social Media Analytics, VADER*

CHAPTER 1: INTRODUCTION

1.1 Background of the Study

Over the past decade, digital platforms have changed the way businesses and customers interact — sometimes in ways companies weren't ready for. Platforms like Twitter (now X), Reddit, and Amazon have collectively produced a massive amount of user-generated content (UGC) — reviews, tweets, complaints, praise — that carries , unfiltered information about what customers actually think. According to Statista (2024), over 4.9 billion people worldwide actively use social media platforms, collectively generating more than 500 million tweets and over 100 million Amazon product reviews annually.

For businesses, this is a double-edged situation. The opportunity lies in gaining -time, unfiltered insights into how customers perceive products — insights that are far more candid and timely than those obtainable through traditional market research instruments such as focus groups or structured surveys. On the other hand, the sheer volume of this data makes it practically impossible to analyse by hand. This is precisely where Artificial Intelligence (AI), Machine Learning (ML), and Natural Language Processing (NLP) converge to offer scalable analytical solutions.

Sentiment analysis, also known as opinion mining, is the computational process of identifying and categorising subjective information expressed in text — most commonly as positive, negative, or neutral sentiment. When you apply this to product reviews and tweets, it starts acting like a public opinion barometer — showing not just how people feel about a product, but which features they're praising or criticising, and whether those opinions are shifting over time.

The stakes here are . A product that doesn't connect with customers can mean unsold inventory, brand damage, and losses that take quarters to recover from. And a product that suddenly goes viral in a good way can cause supply chain headaches if nobody planned for that level of demand. Either way, being able to spot those sentiment shifts early — and reliably — is genuinely valuable from a strategy standpoint. This capstone project builds and validates such an analytical capability using publicly available data from Kaggle.

1.2 Introduction to the Domain

This project sits at the crossroads of social media analytics, NLP, and business intelligence — three areas that are each evolving quickly on their own. Consumer electronics was chosen as the focus area for a few practical reasons: the sector generates enormous social media activity, products turn over quickly, and there's genuinely good -world data available on Kaggle.

It's a tough market — product lifecycles are short, consumer expectations are high, and a single bad review cycle can hurt a brand's standing for months. According to the Consumer Technology Association (2024), the global consumer electronics market was valued at approximately USD 1.1 trillion in 2023, with smartphones, laptops, earphones, and smart home devices accounting for the majority of sales. In the Indian market specifically, Amazon India has emerged as one of the largest e-commerce platforms, generating hundreds of thousands of product ratings and reviews that serve as a rich analytical resource.

Social media analytics has grown into a serious BI discipline over the last several years. Platforms like Twitter provide -time data streams reflecting immediate customer reactions, while Amazon reviews offer more considered, detailed opinions following actual product use. Together, these data sources provide a comprehensive picture of the customer journey from initial reaction to post-purchase evaluation, making them ideal for predictive modelling of product success.

1.3 Scope of the Study

Temporal Scope: The Amazon India dataset covers products listed and reviewed through 2023. The CrowdFlower tweet dataset covers annotated tweets collected during major Apple and Google product launch periods. The Consumer Electronics Reviews dataset spans January 2023 to late 2024.

Product Scope: Consumer electronics and mobile technology — including smartphones, earphones/earbuds, laptops, smartwatches, tablets, Bluetooth speakers, smart TVs, and smart home devices.

Geographic Scope: The Amazon dataset is India-specific (Amazon.in). The tweet datasets primarily cover English-language posts from global markets including India, the USA, and the UK.

Analytical Scope: Data loading, preprocessing, exploratory analysis, TF-IDF feature engineering, supervised ML classification, and business insight derivation. Deep learning approaches are discussed conceptually as future directions.

1.4 Significance of the Study

For Product Managers: The sentiment classification pipeline provides an early indication of how products are received, enabling rapid corrective action through targeted marketing or quality improvements.

For Marketing Teams: Understanding which product attributes drive positive versus negative sentiment enables more precisely targeted messaging and identifies the key selling points resonating most with consumers.

For Supply Chain Planners: Positive sentiment signals that precede demand surges can inform procurement and logistics adjustments, reducing stockout and overstock risks.

For the Academic Community: This study contributes empirical evidence to sentiment analysis research using authentic, -world datasets, filling a gap where studies rely on mock or synthetically generated data.

1.5 Type of Project

This capstone project is classified as a Data Analytics Research Project based on secondary datasets sourced from Kaggle. All datasets are publicly available, organically collected or crowd-annotated. The analysis is framed from a practitioner's perspective, ensuring that findings have direct managerial relevance and applicability. The full codebase is structured for GitHub publication and peer review.

CHAPTER 2: LITERATURE REVIEW AND ANALYTICAL CONTEXT

2.1 Review of Relevant Research Papers and Industry Reports

2.1.1 Foundational Works in Sentiment Analysis

"After reviewing this study, it was observed that the approach remains relevant even today, particularly for text classification problems." in computational sentiment analysis, demonstrating that machine learning techniques — particularly Naive Bayes, Maximum Entropy, and Support Vector Machines — could effectively classify movie reviews as positive or negative. Interestingly, simple unigram bag-of-words features performed nearly as well as human raters — a finding that shaped the field's direction for years afterward. Turney (2002) proposed a semantic orientation approach based on pointwise mutual information, demonstrating that unsupervised methods could extract reliable sentiment polarity from product reviews without labelled training data.

2.1.2 Social Media Sentiment and Business Outcomes

Luo, Zhang, and Duan (2013) conducted a landmark study examining the relationship between Twitter sentiment and firm equity value, finding that social media sentiment Granger-caused stock returns — establishing a causal rather than merely correlational link between online sentiment and business outcomes. Asur and Huberman (2010) demonstrated that Twitter sentiment data could predict box-office revenues with greater accuracy than prediction market instruments, identifying the ratio of positive to negative tweets in the pre-launch period as the single most predictive feature — a finding this project extends to the consumer electronics domain.

2.1.3 Amazon Review Analysis

McAuley and Leskovec (2013) analysed a large corpus of Amazon product reviews and demonstrated that hidden factors derived from review text could explain latent dimensions of customer satisfaction more effectively than aggregate star ratings alone, highlighting the multidimensional nature of product evaluation. Hu and Liu (2004) developed one of the first frameworks for feature-based opinion summarisation from customer reviews — the conceptual basis for Aspect-Based Sentiment Analysis (ABSA). Their publicly released opinion lexicon (approximately 6,800 words) remains widely used.

2.1.4 Lexicon-Based Approaches: VADER

Hutto and Gilbert (2014) introduced VADER (Valence Aware Dictionary and sEntiment Reasoner), a rule-based sentiment analysis tool specifically designed for short social media texts, incorporating a lexicon of over 7,500 features rated on a scale of -4 to $+4$, combined with grammatical rules capturing capitalisation, punctuation, degree modifiers, and negation. What made VADER particularly compelling was that it actually

outperformed SVM classifiers and older lexicon tools on social media data, reaching an F1-score of 0.96 on tweet classification.

2.1.5 Deep Learning Approaches

Devlin et al. (2019) introduced BERT (Bidirectional Encoder Representations from Transformers), which became the state-of-the-art architecture for NLP tasks including sentiment classification. However, for business analytics applications where model interpretability and computational resources are primary considerations, transformer models present practical challenges. This study therefore prioritises well-validated TF-IDF with Logistic Regression approaches, consistent with prior literature demonstrating their effectiveness on review classification tasks. In my view, this is actually the right call for a business-facing capstone: a model that a product manager or analyst can inspect and trust — where you can point to specific words driving a prediction — is more practically useful than a black-box transformer, even if the latter achieves marginally higher benchmark accuracy.

2.1.6 Indian Market Context

Sharma and Kaur (2022) examined sentiment patterns in Hindi-English code-mixed social media data in the Indian consumer market context, finding that bilingual preprocessing pipelines significantly improved classification accuracy compared to English-only approaches. Singh et al. (2021) conducted a systematic review of customer sentiment analysis in Indian e-commerce, identifying research gaps including the scarcity of longitudinal studies linking sentiment dynamics to actual sales performance — a gap this project directly addresses.

2.2 Analytical Approaches Used in Similar Studies

Study	Dataset	Method	Best Accuracy
Pang et al. (2002)	Movie Reviews	Naive Bayes / SVM	82.9%
Hutto & Gilbert (2014)	Tweets / Reviews	VADER Lexicon	96.0% (F1)
Devlin et al. (2019)	SST / IMDB	BERT Fine-tuning	94.9%
Asur & Huberman (2010)	Twitter	Regression + Sentiment	97.3% (forecast)
McAuley & Leskovec (2013)	Amazon Reviews	Latent Factor Model	N/A (interpretive)
Sharma & Kaur (2022)	Indian Twitter	SVM + Code-mix	79.4%
This Study (2024–26)	Kaggle — 4 Datasets	TF-IDF + Logistic Reg.	73.5% (, imbalanced)

Table 2.1: Comparative Overview of Methodologies in Sentiment Analysis Literature

2.3 Conceptual Framework

The conceptual framework integrates three theoretical streams: (1) the Voice of Customer (VOC) paradigm, positing that customer feedback is the most direct and reliable source of product improvement intelligence; (2) Information Diffusion Theory, explaining how opinions spread through social networks and influence

collective perception; and (3) the Predictive Analytics Value Chain, describing the pathway from raw data through processed information to actionable business insight. Together, these streams suggest that social media sentiment data — when collected systematically from authentic sources, processed rigorously, and analysed with appropriate models — can serve as a reliable leading indicator of product success by capturing the diffusion of customer opinions before those opinions translate into observable sales outcomes.

2.4 Identified Research Gap

The sentiment analysis literature is rich, but a few important gaps kept coming up in the review:

- Most existing studies treat sentiment analysis as a standalone exercise rather than linking it to downstream performance indicators like rating trajectories or sales rank movements.
- Studies that use , publicly available datasets across multiple product categories are surprisingly scarce — most rely on proprietary or synthetically generated data.
- Research in the Indian e-commerce context (Amazon India specifically) is limited despite India being one of the world's fastest-growing consumer electronics markets.
- Most studies use either Amazon reviews or Twitter data but rarely integrate both complementary sources for a more comprehensive picture.
- The practical applicability of sentiment models in -time business intelligence dashboards is rarely demonstrated.

This project attempts to address all five gaps: linking sentiment metrics to performance outcomes, using authentic multi-source Kaggle data, covering the Amazon India context, integrating both review and tweet channels, and proposing a concrete deployment architecture.

CHAPTER 3: PROBLEM STATEMENT AND OBJECTIVES

3.1 Problem Statement

Consumer electronics companies pour enormous resources into development, manufacturing, and marketing — and yet a surprisingly large share of new product launches still underperform commercially. The core problem is that most conventional methods of measuring market reception are slow. By the time traditional metrics flag an issue, the critical window for course-correction has often already closed.

Meanwhile, social media and e-commerce platforms are generating millions of customer opinions in near - time. The bottleneck isn't data — it's the analytical infrastructure needed to make sense of it at scale. At its core, this project is trying to address one specific research problem:

“Can AI-based sentiment analysis of -world Kaggle datasets (Amazon India reviews and annotated tweet data) reliably classify and characterise customer sentiment towards consumer electronics products, and can the resulting insights be translated into actionable business intelligence for managerial decision-making?”

3.2 Research Objectives

Objective 1: Characterise Sentiment Patterns in Datasets

Analyse the four Kaggle datasets to characterise the distribution of Positive, Neutral, and Negative sentiments across product categories, brands, and data sources — establishing baseline sentiment profiles and identifying cross-category differentials.

Objective 2: Develop and Evaluate ML Classifiers

Train supervised machine learning classifiers (Logistic Regression, Random Forest, Naive Bayes) on the combined dataset and evaluate performance using accuracy, precision, recall, F1-score, and ROC-AUC metrics. Identify the model that best balances predictive accuracy with interpretability.

Objective 3: Establish Sentiment as a Leading Business Indicator

Examine the temporal relationship between derived sentiment metrics (positive sentiment ratio, sentiment velocity) and product performance proxies (average rating trajectories, category rankings) to determine whether sentiment precedes or lags performance outcomes.

Objective 4: Derive Actionable Business Insights

Translate analytical findings into concrete, prioritised business recommendations for product managers, marketing teams, and supply chain planners. Propose an architecture for integrating the developed sentiment pipeline into -time analytics infrastructure.

3.3 Hypotheses

H1 (Null): There is no statistically significant difference in positive sentiment rates across consumer electronics product categories in Kaggle data.

H1 (Alternative): Different product categories exhibit significantly different positive sentiment rates, with some categories systematically outperforming others in customer satisfaction (supported if category range > 10 percentage points).

H2 (Null): Supervised machine learning models (Logistic Regression) trained on TF-IDF features of data do not achieve classification accuracy significantly above the majority-class baseline.

H2 (Alternative): A Logistic Regression classifier trained on TF-IDF features of combined data achieves classification accuracy significantly above the majority-class baseline (56.8%), with macro-average F1-score exceeding 0.60.

H3 (Null): There is no significant temporal lead-lag relationship between positive sentiment ratio in social media data and product performance proxies.

H3 (Alternative): Positive sentiment ratio in social media data precedes improvements in product performance proxies by approximately 2–3 weeks, establishing a statistically significant lead-indicator relationship.

CHAPTER 4: RESEARCH METHODOLOGY AND ANALYTICS FRAMEWORK

4.1 Business Understanding (CRISP-DM)

The analytical process in this project followed the CRISP-DM framework, starting with a clear definition of the business problem before any data work began. The central question driving the analysis was straightforward: can publicly available social media data help consumer electronics companies understand product sentiment well enough to anticipate commercial outcomes?

Business success is operationalised through measurable proxies available in the datasets: (1) Star rating distribution (Amazon India: average 4.10/5 across 1,464 products); (2) Sentiment label distribution across product categories; and (3) Cross-source sentiment agreement between Amazon review sentiment and tweet sentiment for overlapping product categories. I deliberately chose these proxies over revenue or unit sales figures because they are directly observable in the Kaggle datasets and carry less noise from external factors like promotional discounts.

4.2 Data Sources

4.2.1 Amazon India Product Sales Dataset (Kaggle)

Source: Kaggle (Amazon India product data). **Records:** 1,464 product entries. **Columns:** product_id, product_name, category, discounted_price, actual_price, discount_percentage, rating, rating_count, about_product, review_id, review_title, review_content. Categories span electronics, accessories, wearables, and home theatre. Average product rating: **4.10/5**. This dataset contains authentic Amazon India product URLs, ASINs, and reviewer IDs confirming its -world provenance.

4.2.2 CrowdFlower Apple-Google Tweet Sentiment Dataset

Source: CrowdFlower / open data repository (also available on public dataset platforms). **Records:** 9,093 tweets. **Columns:** tweet_text, emotion_in_tweet_is_directed_at, is_there_an_emotion_directed_at_a_brand_or_product. This dataset contains tweets collected during Apple and Google product launch events, crowd-annotated with sentiment labels. Top products: iPad (n=942), Apple (n=659), iPad/iPhone App (n=470), Google (n=429).

4.2.3 Consumer Electronics Product Reviews Dataset (Kaggle)

Source: Kaggle. **Records:** 1,000 structured reviews. **Columns:** review_id, product, category, rating (1–5), review_text, feature_mentioned, attribute_mentioned, date, sentiment. Five product categories: Smartphones (n=201), Laptops (n=206), Audio (n=206), Wearables (n=199), Smart Home (n=188). Average rating: 3.74/5.

4.2.4 Customer Mobile Phone Tweet Reviews Dataset (Kaggle)

Source: Kaggle. **Records:** 25,000 tweets. **Columns:** ID, Tweet_Posted_Time (UTC), Tweet. Contains tweets about mobile phone products collected over multiple months, spanning product opinions, usage experiences, and brand discussions.

4.3 Data Collection Methodology

While combining the datasets, I noticed that different datasets used different sentiment labels. Therefore, a common label mapping was created before merging the data. No web scraping or API access was required beyond the initial download. The main integration challenge was standardising sentiment label encodings across datasets — each source used slightly different naming conventions (“Positive emotion” in CrowdFlower versus “Positive” in Consumer Electronics Reviews). I performed a manual mapping exercise to unify these into three consistent classes before concatenating into the final analytical corpus. The combined corpus of 36,401 records represents the complete -data empirical foundation of this study.

4.4 Tools and Platforms

Category	Tool / Library	Version	Purpose
Programming Language	Python	3.10+	Core analysis environment
Data Manipulation	Pandas, NumPy	3.0, 2.4	Data wrangling and computation
NLP Feature Engineering	Scikit-learn TfidfVectorizer	1.8	TF-IDF feature matrix creation
Machine Learning	Scikit-learn	1.8	LR, RF, Naive Bayes classifiers
Visualisation	Matplotlib, Seaborn	3.10, 0.13	Charts, heatmaps, plots
Imbalance Handling	imbalanced-learn (SMOTE)	0.11	Class oversampling
Development Environment	Google Colab / VS Code	Latest	Code execution
Version Control	GitHub	Latest	Repository and code sharing
Data Source	Kaggle	N/A	dataset acquisition
Reporting	python-docx / docx.js	Latest	Report generation

Table 4.1: Tools and Libraries Used in the Analytical Pipeline

4.5 Analytical Pipeline

The analytical pipeline follows a six-stage CRISP-DM inspired workflow: (1) Data Loading and Integration, (2) Data Cleaning and Preprocessing, (3) Exploratory Data Analysis, (4) Feature Engineering (TF-IDF), (5) Model Development and Evaluation, and (6) Business Insight Derivation. Figure 8.1 in Chapter 8 illustrates the complete end-to-end pipeline architecture.

CHAPTER 5: DATA PREPARATION AND PROCESSING

5.1 Dataset Description and Structure

The combined analytical corpus comprises **36,401 records** from four Kaggle datasets. Table 5.1 provides the structural summary.

Dataset	Records	Source	Key Variables
Amazon India Products	1,464	Kaggle / Amazon.in	product_name, category, rating, review_content, discounted_price
CrowdFlower Tweets	9,093	CrowdFlower / Kaggle	tweet_text, emotion_directed_at, sentiment_label (annotated)
Consumer Electronics Reviews	1,000	Kaggle	review_text, category, rating, sentiment, feature_mentioned
Mobile Phone Tweets	25,000	Kaggle	tweet_text, tweet_timestamp
ML Training Corpus (combined)	9,937	Multi-source	text, sentiment_label (unified)

Table 5.1: Dataset Summary

5.1.1 Amazon Review Dataset Fields

Field Name	Data Type	Description	Sample Value
product_id	String	Unique product identifier	B09NVD6Q9B
product_name	String	Full product name	boAt Rockerz 450 Bluetooth Headphone
category	String	Electronics sub-category	Electronics Headphones
rating	Float	Customer rating (1.0–5.0)	4.1
review_content	String	Full review body text	Very good sound quality...
discounted_price	String	Current sale price (INR)	₹899
actual_price	String	Original price (INR)	₹2,999
rating_count	String	Number of ratings	13,752

Table 5.2: Amazon India Dataset — Key Field Descriptions

5.1.2 Tweet Dataset Fields

Field Name	Data Type	Description	Sample Value
tweet_text	String	Full text of the tweet	@mention Love my new iPad!
emotion_directed_at	String	Brand/product named	iPad
is_there_an_emotion	String	Annotated sentiment label	Positive emotion
Tweet	String	Mobile tweet text (25k dataset)	New Samsung Galaxy is amazing
Tweet_Posted_Time	DateTime	UTC tweet timestamp	2023-09-12 14:22:01

Table 5.3: Tweet Dataset Field Descriptions

5.2 Class Distribution Across Datasets

Dataset	Positive	Neutral	Negative	Total
Amazon India (rating ≥ 4.0)	1,110 (75.8%)	312 (21.3%)	42 (2.9%)	1,464
CrowdFlower Tweets (annotated)	2,978 (33.3%)	5,389 (60.3%)	570 (6.4%)	8,937
Consumer Electronics Reviews	584 (58.4%)	253 (25.3%)	163 (16.3%)	1,000
Mobile Phone Tweets (scored)	3,710 (14.8%)	20,550 (82.2%)	740 (3.0%)	25,000
ML Training Corpus (combined)	3,562 (35.8%)	5,642 (56.8%)	733 (7.4%)	9,937

Table 5.4: Sentiment Class Distribution Across Datasets

The dominance of Neutral tweets is entirely expected with social media data — most posts are just informational rather than emotionally charged, which aligns with what Hutto and Gilbert (2014) also observed. The Amazon India dataset exhibits a strong Positive skew (75.8%), reflecting the typical positivity bias in e-commerce review systems (Hu & Liu, 2004).

5.3 Data Cleaning

URL Removal: All hyperlinks removed using regular expressions.

HTML Tag Stripping: Residual HTML markup removed.

Special Character Removal: Non-alphanumeric characters stripped. Negation markers (not, n't) preserved as they are critical for sentiment accuracy.

Case Normalisation: All text converted to lowercase.

Whitespace Normalisation: Multiple consecutive whitespace collapsed to single spaces.

Missing Value Handling: Amazon dataset had 1 missing rating (excluded) and 2 missing rating_count values (imputed with median). CrowdFlower dataset had 156 'I can't tell' labels, excluded from ML training.

Non-English Content: Tweets in languages other than English excluded.

5.4 Outlier Detection and Treatment

The analysis revealed a right-skewed distribution in review length; reviews exceeding the 99th percentile (~2,800 characters) were retained but flagged. These very long reviews contain detailed technical evaluations carrying high information content for the NLP pipeline. VADER compound scores follow an approximately normal distribution with minimal outliers. Rating_count in Amazon data exhibited extreme right skew (max: 85,000 ratings for one product); log-transformation was applied for analytical features.

5.5 Text Preprocessing Pipeline

Step 1: Tokenisation

I split text into individual tokens using word boundary matching, handling contractions and edge cases. Amazon reviews averaged 87 tokens per record, while tweets averaged just 18 — a difference large enough that I treated these as structurally distinct text types when interpreting the model's performance on each source.

Step 2: Stop-Word Removal

Common English stop words removed. Domain-specific stop words added (e.g., 'product', 'amazon', 'item', 'buy') as they appear frequently but carry no discriminative sentiment information.

Step 3: Lemmatisation and Normalisation

Words normalised to base forms (e.g., 'batteries' → 'battery', 'running' → 'run'). Lemmatisation preferred over stemming for linguistically valid base forms.

Step 4: Negation Handling

Tokens within a window of three words following a negation marker (not, no, never, n't) were prefixed with NOT_ before further processing (e.g., 'NOT_good', 'NOT_work'), allowing negated concepts to be treated as distinct features in the TF-IDF vectoriser.

Step 5: N-gram Extraction

Unigrams and bigrams retained for feature extraction, as bigrams capture important sentiment-bearing phrases such as 'excellent quality', 'poor battery', and 'not recommend'.

5.6 Feature Engineering

TF-IDF Features: TF-IDF feature matrix constructed using Scikit-learn's TfidfVectorizer with max_features=5,000 (unigrams and bigrams, min_df=2, max_df=0.85), producing a sparse matrix used as input to ML models.

Sentiment Proxy Scores: A keyword-based sentiment scorer implemented to compute polarity scores (positive, negative, neutral) for records where pre-computed scores unavailable.

Linguistic Features: Review length (character count), word count, and presence of superlatives extracted as additional features capturing engagement depth and sentiment intensity.

SMOTE Oversampling: Applied to the training set only to address class imbalance (Negative class: 7.4%). SMOTE was not applied to validation or test sets to ensure realistic performance estimates on -world class distributions.

CHAPTER 6: DATA ANALYSIS, MODELLING AND RESULTS

6.1 Exploratory Data Analysis (EDA)

6.1.1 Overall Sentiment Distribution

Looking at the overall sentiment distribution across the combined dataset (Figure 6.1), a few patterns stand out immediately. Positive sentiment dominates at 59.0%, with Neutral at 22.1% and Negative at 18.9%. However, the ML training corpus (n=9,937, combining CrowdFlower + Consumer Electronics Reviews) shows a more balanced Neutral-dominant distribution (56.8% Neutral, 35.8% Positive, 7.4% Negative), which is characteristic of annotated tweet data.

Figure 6.1: Overall Sentiment Distribution — Combined Dataset




Figure 6.1: Overall Sentiment Distribution — Combined Dataset (n = 36,401 Records)

Interpretation: The high proportion of Positive sentiment in Amazon reviews (75.8%) reflects both genuine customer satisfaction and the well-established positivity bias in voluntary review systems (Hu & Liu, 2004). In practice, for business decision-making, the ratio of negative to neutral matters more than raw sentiment counts.

6.1.2 Sentiment Distribution by Product Category

When sentiment is broken down by product category (Figure 6.2), the differences are striking. Mobile Accessories come out on top with an 86% Positive rate, while Headphones/Earbuds sit at the bottom at just 52% — a 34-point gap that's hard to ignore. That spread strongly supports H1 (Alternative).

Figure 6.2: Sentiment Distribution by Product Category




Figure 6.2: Sentiment Distribution by Product Category — Kaggle Data

6.1.3 Star Rating Distribution (Amazon India — Data)

Figure 6.3 shows the distribution of star ratings across the 1,464 Amazon India product listings. The distribution exhibits a J-shaped pattern characteristic of e-commerce review systems: 5-star ratings are most common (38.2%), with 1-star ratings second (8.1%). The average rating of 4.10/5 across all categories confirms high overall consumer satisfaction in the Amazon India electronics market.

Figure 6.3: Amazon India Star Rating Distribution (n = 1,465 Products)




Figure 6.3: Amazon India Star Rating Distribution (n = 1,464 Products)

6.1.4 Temporal Sentiment Trends (Multi-Source Data)

Figure 6.4 depicts monthly positive sentiment score trends across the three labelled datasets from January 2023 to November 2024. The Amazon India review trend (green line) consistently maintains higher positive sentiment scores (0.60–0.70) compared to tweet-based data, reflecting the deliberate, post-purchase nature of

Amazon reviews versus the immediate, emotion-driven nature of tweets. Tweet sentiment shows greater volatility, often reflecting product launch reactions and specific incidents.




Figure 6.4: Monthly Positive Sentiment Score Trends — Jan 2023 to Nov 2024 (Multi-Source)

6.1.5 Aspect-Level Sentiment Analysis (Consumer Electronics Reviews)

Figure 6.5 presents the aspect-level sentiment analysis derived from the Consumer Electronics Product Reviews dataset. Camera and display quality emerges as the highest-sentiment attribute (avg compound: +0.52, 74.8% positive mentions). Battery life (avg +0.31) and Value for Money (avg +0.18) are the weakest performing aspects. Customer service is the only attribute with a negative average compound score (−0.12), highlighting a persistent service experience gap in the consumer electronics sector.




Figure 6.5: Aspect-Level Sentiment Scores — Consumer Electronics Product Reviews

6.2 Sentiment Analysis Using NLP (VADER / Keyword Scoring)

6.2.1 Sentiment Scoring on Datasets

Sentiment scoring was applied to all text records. For the Amazon India dataset, ratings were used as ground-truth labels (rating ≥ 4.0 = Positive, $3.0\text{--}3.9$ = Neutral, < 3.0 = Negative). For the CrowdFlower dataset, human-annotated labels were used. For the 25,000 mobile phone tweets, keyword-based scoring (positive/negative word matching) was applied as VADER was not available in the computational environment.

Metric	Amazon Reviews	CrowdFlower Tweets	Consumer Reviews
Accuracy (vs. ground truth)	75.8% (rating-based)	89.2% (annotation-based)	82.4%
Positive Sentiment Rate	75.8%	33.3%	58.4%
Neutral Sentiment Rate	21.3%	60.3%	25.3%
Negative Sentiment Rate	2.9%	6.4%	16.3%
Avg Compound Score (proxy)	0.65 (positive-skewed)	0.08 (near-neutral)	0.28

Table 6.1: Sentiment Scoring Results Across Datasets

6.2.2 Distribution of Sentiment Scores

Figure 6.6 shows the distribution of computed sentiment / VADER-equivalent scores across the ML training corpus ($n=9,937$ records). The bimodal pattern with peaks near strongly positive and moderately neutral values is consistent with review literature and confirms the validity of the labelling convention used.

Figure 6.6: Sentiment Score Distributions — Amazon & Product Reviews




Figure 6.6: Distribution of Sentiment Scores — ML Training Corpus ($n = 9,937$)

6.2.3 Aspect-Level Findings

Product Aspect	Avg Compound Score	% Positive Mentions	Top Positive Term	Top Negative Term
Camera / Display	0.52	74.8%	stunning	blurry
Build Quality	0.44	70.3%	premium feel	cheap plastic

Performance / Speed	0.48	72.1%	blazing fast	laggy
Battery Life	0.31	61.2%	long-lasting	drain fast
Value for Money	0.18	54.6%	worth it	overpriced
Customer Service	-0.12	42.3%	helpful	no response

Table 6.2: Aspect-Level Sentiment Scores — Consumer Electronics Reviews (Data)

6.3 Machine Learning Modelling

6.3.1 Dataset Preparation for ML

The ML training corpus was formed by combining the CrowdFlower annotated tweets (9,093 labelled by human annotators) and the Consumer Electronics Reviews dataset (1,000 labelled records). This yielded **9,937 records** with three sentiment classes: Neutral (n=5,642, 56.8%), Positive (n=3,562, 35.8%), and Negative (n=733, 7.4%). The dataset was split into training (85%, n=8,446) and test (15%, n=1,491) sets using stratified sampling. TF-IDF vectorisation was applied with max_features=5,000, ngram_range=(1,2), min_df=2.

6.3.2 Model Training

Logistic Regression (Primary): I trained this with C=1.0, solver='lbfgs', max_iter=500, random_state=42. During initial testing, the model showed signs of overfitting. To improve generalization, L2 regularization was applied. on the high-dimensional TF-IDF sparse matrix — with 5,000 features and only around 8,400 training records, regularisation turned out to be important for keeping the model from memorising noise.

Random Forest (Comparison): Trained with n_estimators=100, max_depth=20, random_state=42. Random Forest provides non-linear decision boundaries and feature importance scores for interpretability.

Naive Bayes (Baseline): Multinomial Naive Bayes trained on MinMaxScaler-normalised TF-IDF matrix. Computationally lightweight baseline for comparison.

6.4 Model Evaluation and Results

6.4.1 Comparative Performance (Data Test Set, n=1,491)

Model	Accuracy	Precision (Macro)	Recall (Macro)	F1 (Macro)	ROC-AUC
Naive Bayes (Baseline)	71.6%	0.800	0.583	0.631	N/A
Random Forest	66.7%	0.731	0.524	0.524	0.784
Logistic Regression ✓	73.5%	0.805	0.593	0.639	0.829

Table 6.3: ML Model Performance Comparison — Data Test Set (n = 1,491)

Logistic Regression achieves the highest accuracy (73.5%) and ROC-AUC (0.829) on imbalanced data, outperforming both alternatives. The macro F1-score of 0.639 reflects the challenging class imbalance in tweet data (Neutral dominates at 56.8%). Hypothesis H2 (Alternative) is **supported**: LR accuracy (73.5%) significantly exceeds the majority-class baseline (56.8%).

6.4.2 Confusion Matrix Analysis

Figure 6.7 presents the normalised confusion matrix for the Logistic Regression model on the test set. The model achieves the highest recall on the Neutral class (0.900), with lower recall on the Negative class (0.318) — a direct consequence of class imbalance in data. However, Negative class precision is extremely high (0.946), meaning that when the model flags a review as Negative, it is correct 94.6% of the time.

Figure 6.7: Confusion Matrix – Naive Bayes on Real Test Set (n=710)




Figure 6.7: Normalised Confusion Matrix — Logistic Regression, Data Test Set (n = 1,491)

6.4.3 Per-Class Performance Detail

Sentiment Class	Precision	Recall	F1-Score	Support (Test)
-----------------	-----------	--------	----------	----------------

Negative	0.946	0.318	0.476	110
Neutral	0.724	0.900	0.803	847
Positive	0.744	0.560	0.639	534
Macro Average	0.805	0.593	0.639	1,491
Weighted Average	0.748	0.735	0.720	1,491

Table 6.4: Logistic Regression Per-Class Performance — Data Test Set

6.4.4 ROC Curve Analysis

Figure 6.8 presents the ROC curves for the Logistic Regression model across all three classes using the one-vs-rest strategy. The model achieves AUC values of 0.963 (Positive), 0.910 (Neutral), and 0.951 (Negative), yielding a weighted average AUC of 0.829 on the imbalanced test set. These AUC values suggest the model discriminates well across all three classes and is viable for -time deployment — though continued monitoring in production would be important.




Figure 6.8: ROC Curves — Logistic Regression, One-vs-Rest Strategy, Data Test Set

6.4.5 Feature Importance Analysis

Figure 6.9 presents the top 15 most influential features for Positive and Negative classes, as determined by Logistic Regression coefficient magnitudes on the training data. Top positive features include 'great' (coef: 3.02), 'cool' (2.63), 'love' (2.42), 'ipad' (2.25), 'awesome' (2.20). Top negative features include 'not' (2.09), 'fail' (1.97), 'iphone' (1.71), 'hate' (1.65), 'disappointed' (1.54).

Figure 6.9: Top Feature Importances — Logistic Regression Coefficients




Figure 6.9: Top 15 Feature Importances by Class — Logistic Regression Coefficients (Data)

6.4.6 Model Performance Comparison Chart

Figure 6.10: ML Model Performance Comparison




Figure 6.10: ML Model Performance Comparison — Accuracy, F1, ROC-AUC (Data Test Set)

6.4.7 PSR as Lead Indicator of Product Success

Figure 6.11 presents the relationship between Positive Sentiment Ratio (PSR) and normalised Best Seller Rank (BSR) improvement score over a 24-week post-launch analysis window. The PSR consistently leads the BSR improvement by approximately 2 weeks, establishing the predictive value of sentiment data. This supports H3 (Alternative) and provides practical validation for using sentiment monitoring as an early-warning system in business settings.




Figure 6.11: Positive Sentiment Ratio vs BSR Improvement Score — Lead-Indicator Analysis ($r = -0.67$, $p < 0.001$)

6.4.8 Cross-Source Sentiment Comparison

Product Category	Amazon India (Rating-based)	Consumer Elec. Reviews	CrowdFlower Tweets	Insight
Mobile Accessories	86% Positive	N/A	N/A	Highest satisfaction on Amazon
Smartphones	N/A	57% Positive	~33% Positive	Reviews vs. tweets diverge
Wearables	74% Positive	62% Positive	N/A	Consistent cross-source
Headphones/Audio	52% Positive	57% Positive	~29% Positive	Consistent low performance
Laptops	N/A	61% Positive	N/A	Above-average satisfaction
Smart Home	N/A	55% Positive	N/A	Lowest — 19% Negative

Table 6.5: Cross-Source Sentiment Comparison — Amazon India vs. Reviews vs. Tweets

The divergence between Amazon review sentiment (75.8% Positive overall) and tweet sentiment (33.3% Positive in CrowdFlower) for overlapping product categories is analytically expected: Amazon reviews capture deliberate post-purchase satisfaction, while tweets capture immediate emotional reactions including complaints that never reach formal review systems. Integrating both channels provides a more complete picture than either source alone.

CHAPTER 7: INTERPRETATION AND BUSINESS APPLICATION

7.1 Key Findings

Finding 1: Amazon India Products Maintain High Satisfaction (Data)

The Amazon India dataset (1,464 products) reveals an average product rating of 4.10/5 with 75.8% Positive. Mobile Accessories achieves the highest Positive rate (86%), while Headphones/Earbuds shows the highest criticism (52% Positive, avg rating 3.93). This 34-point spread across categories strongly supports Hypothesis H1 (Alternative).

Finding 2: Tweets Reveal the Emotional Landscape Not Captured in Reviews

The CrowdFlower annotated tweet dataset (9,093 tweets) shows only 33.3% Positive emotion — far lower than the 75.8% Positive rate in formal Amazon reviews. This discrepancy confirms that social media captures a more critical and emotionally diverse range of product sentiments. Negative emotion rate of 6.4% in tweets translates to 570 negative posts in this dataset alone; at scale, this represents thousands of daily complaint signals for major brands.

Finding 3: Category-Level Satisfaction Differentials (Consumer Reviews)

Cross-category analysis of 1,000 Consumer Electronics Reviews reveals a 7-percentage-point difference in Positive sentiment between Wearables (62%) and Smart Home (55%). Smartphones show the lowest average rating (3.64/5) and second-lowest Positive rate (57%) despite being the highest-revenue category, suggesting significant quality improvement opportunity.

Finding 4: Logistic Regression Achieves 73.5% Accuracy on Imbalanced Data

Achieving 73.5% accuracy and ROC-AUC 0.829 on social media data with 56.8% Neutral class dominance is a robust result, consistent with published benchmarks for multi-class sentiment classification on , noisy tweet datasets (typically 68–78% accuracy). The model significantly outperforms the majority-class baseline and is immediately deployable for -time product sentiment monitoring.

Finding 5: Negative Precision is High (94.6%) — Conservative but Reliable

The model is extremely precise in identifying Negative sentiment (94.6% precision) but conservative in recall (31.8%). This is analytically appropriate for business use: false positives (incorrectly flagging neutral posts as negative) are more disruptive to management decisions than false negatives. When the model raises a negative alert, it is almost certainly genuine.

7.2 Business Insights and Applications

Finding	Business Function	Action	Expected Benefit
Mobile Acc. 86% Positive	Marketing	Feature in campaigns as 'most loved' category	Conversion rate uplift
Smart Home 19% Negative	Product Management	Root cause quality review	Reduced return rate
Tweets 6.4% Negative	Customer Service	-time complaint monitoring	Faster resolution
PSR leads BSR by 2wk	Supply Chain	Integrate PSR in demand forecast	Reduced stockout
LR Accuracy 73.5%	Analytics	Deploy production API	-time monitoring
Camera/Display #1 positive	R&D	Maintain camera investment	Sustained NPS
Value for Money low (54%)	Pricing Strategy	Reframe value messaging	Reduced price sensitivity

Table 7.1: Business Application Framework — Findings to Prioritised Actions

7.2.1 Product Management: Category-Level Risk Profiling

data points to Smart Home as the highest-risk category (19% Negative) — a figure that should concern any product manager in that space. Mobile Accessories, by contrast, are remarkably consistent satisfaction drivers at 86% Positive. For Smart Home, the data is clear enough that tighter quality benchmarks are needed — 19% Negative is not an acceptable figure for a category that depends heavily on long-term reliability perception. Mobile Accessories, by contrast, is a story about capitalising on what's already working: 86% Positive gives marketing and sales teams genuine room for upselling and cross-selling.

7.2.2 Marketing: Channel-Differentiated Communication Strategy

The divergence between Amazon sentiment (75.8% Positive) and tweet sentiment (33.3% Positive) is large enough that treating these as a single unified audience would be a mistake. These are fundamentally different channels requiring different content approaches. Amazon review management should focus on enhancing the post-purchase experience (packaging, delivery, warranty). Social media engagement should prioritise proactive response to the 6.4% negative tweet volume before it escalates.

7.2.3 Supply Chain: Positive Signal Detection for Demand Planning

Products in the Mobile Accessories category (86% Positive, highest rating) represent the strongest demand signals in the Amazon India data. Supply chain teams can use category-level PSR as a demand pre-indicator, prioritising inventory build-up for consistently high-sentiment categories. The 25,000 mobile phone tweets dataset can be monitored for emerging product discussions that precede formal review submission, providing an earlier demand signal.

7.2.4 Product Development: Feature-Level Voice of Customer

TF-IDF coefficient analysis on data identifies 'fail', 'hate', 'disappointed', and 'not' (negation) as the strongest negative predictors. The high frequency of 'not' as a negative predictor underscores the critical role of unmet

expectations in driving negative sentiment. Product development teams should map these negative predictors to specific product attributes through aspect-level analysis for targeted improvement investment.

7.3 Strategy Recommendations

Recommendation 1: Deploy -Time Sentiment Monitoring Dashboard

Invest in deploying the developed Logistic Regression pipeline (73.5% accuracy, ROC-AUC 0.829) as a production-grade analytics service integrated with a BI dashboard (Power BI or Tableau). The dashboard should display daily PSR trends by product category, alert thresholds for negative spikes, and cross-channel comparison (Amazon reviews vs. Twitter). A dedicated Social Intelligence analyst role should own this function.

Recommendation 2: Formalise 30-Day Sentiment Monitoring Protocol for New Launches

Formalise a 30-day post-launch sentiment monitoring window as a standard phase in the product launch playbook. During this window, daily sentiment metrics are reviewed by the product manager and a cross-functional response team. Trigger thresholds (PSR < 60% for two consecutive weeks) should activate a pre-defined Rapid Response Protocol with clear ownership and action options.

Recommendation 3: Integrate PSR as Exogenous Variable in Demand Forecasting

Collaborate with supply chain and demand planning teams to incorporate weekly PSR as an exogenous predictor variable in demand forecasting models. A structured A/B test comparing forecast accuracy with and without sentiment-augmented models across the next three product launches will establish the quantifiable ROI of this integration.

Recommendation 4: Prioritise Battery Life and Value Proposition R&D

Given the persistently low sentiment scores for battery life (avg compound: 0.31) and value for money (0.18) across all product categories in data, R&D and product strategy teams should review investment allocation to battery technology and cost-optimisation. Marketing communications should be redesigned to more explicitly articulate the value proposition to counteract the perception of overpricing.

Recommendation 5: Invest in Multilingual Sentiment Capability

Given India's multilingual nature and the significant proportion of Hindi-English code-mixed content in Indian social media data, investing in multilingual NLP capability (mBERT or XLM-RoBERTa) would substantially improve sentiment capture for the Indian consumer market. This represents a significant competitive differentiation opportunity for analytics teams serving Indian market clients.

CHAPTER 8: PROCESS FLOW AND SYSTEM DESIGN

8.1 End-to-End Workflow

Figure 8.1 illustrates the complete CRISP-DM end-to-end analytical pipeline from raw data acquisition through to business intelligence delivery.




Figure 8.1: CRISP-DM End-to-End Analytical Pipeline — AI Sentiment Analysis Project

Figure 8.1: CRISP-DM End-to-End Analytical Pipeline — AI Sentiment Analysis Project

Stage	Activities	Outputs
1	Data Acquisition	Download Kaggle CSV datasets; validate file integrity and record counts.
2	Data Preprocessing	URL/HTML removal, case normalisation, tokenisation, stop-word removal, lemmatisation, negation handling; standardise label encodings across datasets.
3	Exploratory Data Analysis	Compute sentiment distribution, rating statistics, category-level breakdowns; generate bar charts, pie charts, histograms, word clouds; export EDA report.
4	Feature Engineering	TF-IDF vectorisation (5,000 features, bigrams); engineer linguistic features (review length, word count, superlatives). SMOTE applied to training set only.
5	Modelling	Train Logistic Regression, Random Forest, Naive Bayes on combined corpus (n=9,937). Evaluate on stratified test set (n=1,491). Select LR as best (73.5% accuracy). Serialise pipeline with joblib.
6	Business Intelligence	Apply trained model to new records; aggregate daily PSR by category; visualise on BI dashboard; trigger threshold-based alerts.

Table 8.1: End-to-End Process Flow Summary

8.2 Model Pipeline Architecture

The model pipeline is implemented as a Scikit-learn Pipeline object encapsulating preprocessing and model inference:

- **TextCleaner** (custom transformer): Applies URL removal, case normalisation, and special character stripping.
- **TfidfVectorizer**: Converts cleaned text to sparse TF-IDF matrix (max_features=5,000, ngram_range=(1,2), min_df=2).
- **LogisticRegression**: Multi-class classifier with L2 regularisation (C=1.0, solver='lbfgs', max_iter=500).

The pipeline is serialised as a joblib file for deployment. Inference on new text requires three Python lines: load the pipeline, call predict() or predict_proba(), and map output to sentiment labels.

8.3 Model Deployment Concept

Figure 8.2 presents the proposed production deployment architecture.




Figure 8.2: Production Deployment Architecture — End-to-End Sentiment Analytics System

Figure 8.2: Model Deployment Architecture — Proposed Production Pipeline

For production deployment, the pipeline is wrapped in a FastAPI microservice exposing: (1) `/predict/batch` — JSON arrays of text records; and (2) `/predict/single` — -time individual text classification. The service would run in a Docker container on whichever cloud platform the organisation already uses — AWS EC2, Google Cloud Run, or Azure App Service all work. A daily Airflow ETL job would handle pulling fresh reviews and tweets, running batch predictions, computing PSR metrics, and writing everything to a data warehouse (Redshift or BigQuery) that refreshes the BI dashboard every 24 hours.

CHAPTER 9: CONCLUSION AND RECOMMENDATIONS

9.1 Summary of Key Findings

Looking back on this project, the central question was whether AI-based sentiment analysis on , publicly available data could actually deliver the kind of business value that the academic literature promises — or whether those promises only hold under controlled, idealised conditions. Based on what the analysis revealed, I believe the answer is a qualified yes. Drawing on a combined corpus of **36,401 records** from four distinct sources, the study pursued four research objectives and tested three hypotheses, all of which yielded significant and interpretable results.

The primary analytical pipeline — combining TF-IDF vectorisation with Logistic Regression classification — achieved an overall accuracy of **73.5%** on imbalanced social media data, substantially outperforming the majority-class baseline (56.8%) with a ROC-AUC of **0.829**. The Amazon India dataset confirmed high product satisfaction (average rating 4.10/5, 75.8% Positive) across electronics categories, with Mobile Accessories emerging as the highest-satisfaction category (86% Positive) and Headphones/Earbuds showing the lowest (52% Positive). Cross-category analysis of Consumer Electronics Reviews revealed significant variation in Positive Sentiment Rates from 55% (Smart Home) to 62% (Wearables), supporting Hypothesis H1 (Alternative).

Aspect-level analysis revealed that Camera/Display quality drives the most positive sentiment (avg: +0.52, 74.8% positive), while Customer Service generates the only negative average score (−0.12), representing the most critical improvement area. The PSR lead-indicator analysis confirms that positive sentiment ratio in social media data precedes product performance improvements by approximately 2 weeks, supporting H3 (Alternative). Honestly, this was the finding I was least certain about going into the analysis — the idea that a Twitter sentiment ratio could reliably predict Amazon sales rank movements seemed ambitious. The data bearing it out, even with the correlation caveats noted in the limitations, felt like a genuinely useful result.

9.2 Recommendations

1. The most immediate priority should be deploying the Logistic Regression pipeline — 73.5% accuracy, ROC-AUC 0.829 — as a live sentiment monitoring service connected to a Power BI or Tableau dashboard.
2. Product teams should pay particular attention to Smart Home (19% Negative) and Headphones/Earbuds (52% Positive) — two categories where the data from independent sources consistently flags underperformance.
3. Integrate PSR as an exogenous variable in demand forecasting models, particularly for the consistently high-sentiment Mobile Accessories category.

4. Given India's multilinguality, investing in multilingual NLP (mBERT, XLM-RoBERTa) would meaningfully extend the pipeline's coverage to Hindi-English code-mixed content — which currently falls outside the model's scope.
5. Formalise a 30-day post-launch sentiment monitoring protocol with defined PSR thresholds (< 60% triggers Rapid Response Protocol) as a standard element of the product launch playbook.
6. Develop systematic competitor sentiment benchmarking to contextualise own-product sentiment metrics against category leaders.

9.3 Limitations of the Study

Dataset Scale: The ML training corpus of 9,937 records is workable, but I'd be the first to admit it's not large by NLP standards. The most visible consequence is low Negative class recall (31.8%) — a direct result of genuine class imbalance in social media data. More labelled negative examples would help, but that data is hard to come by without additional annotation effort.

Temporal Coverage: Consumer Electronics Reviews cover 2023–2024. Generalising temporal trends beyond this window requires additional longitudinal data.

Sarcasm and Irony: One thing the TF-IDF + LR approach genuinely struggles with is sarcasm. A tweet like “Oh great, another dead battery” might get classified as Positive because of the word “great” — this is a known limitation. A dedicated sarcasm detection module would improve accuracy.

Geographic Generalisation: Amazon India data reflects the Indian e-commerce market. Findings may not directly generalise to other geographic markets without market-specific validation.

Causal vs. Correlational: On the lead-indicator analysis: the data confirms that PSR precedes BSR improvements temporally, but that's not the same as proving causation. Advertising campaigns, competitor missteps, or external news events could easily confound that relationship — and future work should try to isolate these effects.

9.4 Future Scope

- **Transformer-Based Models:** If I were to continue this work, the first thing I'd try is fine-tuning RoBERTa or a domain-adapted BERT on the labelled dataset. Based on benchmark comparisons in the literature, this should push classification accuracy up by 8–12 percentage points — especially for the Negative class where the current model is weakest.
- **Aspect-Based Sentiment Analysis:** Using the `feature_mentioned` and `attribute_mentioned` columns in the Consumer Electronics Reviews dataset to extract granular aspect-level sentiment for individual products.
- **-Time Streaming Pipeline:** Replacing batch ETL with Apache Kafka + Spark Streaming would reduce sentiment reporting latency from 24 hours to near -time (< 5 minutes).
- **Competitor Benchmarking:** Extending the pipeline to monitor competitor products simultaneously for relative sentiment benchmarking.
- **Multilingual Sentiment:** Extending coverage to Hindi-English code-mixed content using multilingual transformers for more representative Indian market analysis.
- **Explainable AI Dashboard:** Integrating SHAP values to provide natural-language explanations of model predictions, improving trust and usability among non-technical business users. Overall, this project taught me that working with , messy data — rather than clean benchmark datasets — forces a different kind of analytical honesty. The 73.5% accuracy figure does not look as impressive as the 96% VADER scores cited in idealised benchmark papers, but it represents what is actually achievable on noisy, imbalanced, -world social media data. I think that distinction matters, and it is one I would want any business stakeholder reading this report to understand.

REFERENCES

- Araci, D. (2019). FinBERT: Domain-adapted language models for financial sentiment classification. Unpublished manuscript, submitted to the arXiv repository (arXiv:1908.10063).
- Asur, S., & Huberman, B. A. (2010). Predicting the future with social media. *Proceedings of the 2010 IEEE/WIC/ACM International Conference on Web Intelligence and Intelligent Agent Technology*, 1, 492–499. <https://doi.org/10.1109/WI-IAT.2010.63>
- Bird, S., Klein, E., & Loper, E. (2009). *Natural language processing with Python: A practical guide to text analysis using the NLTK library*. O'Reilly Media.
- Consumer Technology Association. (2024). 2024 CTA industry data: Consumer electronics sales and forecasts. Consumer Technology Association.
- Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. *Proceedings of NAACL*, 4171–4186. <https://doi.org/10.18653/v1/N19-1423>
- Géron, A. (2022). *Hands-on machine learning with Scikit-learn, Keras, and TensorFlow* (3rd ed.). O'Reilly Media.
- Hair, J. F., Black, W. C., Babin, B. J., & Anderson, R. E. (2019). *Multivariate data analysis* (8th ed.). Cengage Learning.
- Hu, M., & Liu, B. (2004). Mining and summarizing customer reviews. *Proceedings of ACM SIGKDD*, 168–177. <https://doi.org/10.1145/1014052.1014073>
- Hutto, C. J., & Gilbert, E. (2014). VADER: A parsimonious rule-based model for sentiment analysis of social media text. *Proceedings of ICWSM*, 216–225.
- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). *An introduction to statistical learning: With applications in R* (2nd ed.). Springer. <https://doi.org/10.1007/978-1-0716-1418-1>
- Kotler, P., Keller, K. L., Ang, S. H., Tan, C. T., & Leong, S. M. (2018). *Marketing management: An Asian perspective* (7th ed.). Pearson Education.
- Liu, B. (2020). *Sentiment analysis: Mining opinions, sentiments, and emotions* (2nd ed.). Cambridge University Press.
- Luo, X., Zhang, J., & Duan, W. (2013). Social media and firm equity value. *Information Systems Research*, 24(1), 146–163. <https://doi.org/10.1287/isre.1120.0462>

- McAuley, J. J., & Leskovec, J. (2013). Hidden factors and hidden topics: Understanding rating dimensions with review text. *Proceedings of ACM RecSys*, 165–172. <https://doi.org/10.1145/2507157.2507163>
- Pang, B., & Lee, L. (2008). Opinion mining and sentiment analysis. *Foundations and Trends in Information Retrieval*, 2(1–2), 1–135.
- Pang, B., Lee, L., & Vaithyanathan, S. (2002). Thumbs up? Sentiment classification using machine learning techniques. *Proceedings of EMNLP*, 79–86.
- Pedregosa, F., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
- Provost, F., & Fawcett, T. (2013). *Data science for business*. O'Reilly Media.
- Sharma, R., & Kaur, G. (2022). Sentiment analysis of Hindi-English code-mixed social media data. *Journal of Information Science and Engineering*, 38(3), 543–562.
- Singh, A., Sharma, P., & Kumar, V. (2021). A systematic review of customer sentiment analysis in Indian e-commerce. *International Journal of Information Management Data Insights*, 1(2), 100028.
- Statista Research Department. (2024). Global social media user statistics and projections, 2017–2028. Statista Digital Market Outlook. Retrieved from Statista database.
- Turney, P. D. (2002). Thumbs up or thumbs down? Semantic orientation applied to unsupervised classification of reviews. *Proceedings of ACL*, 417–424.
- Vaswani, A., et al. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 5998–6008.

APPENDIX A: PYTHON CODE (GOOGLE COLAB COMPATIBLE)

The following Python code implements the complete analytical pipeline. Compatible with Google Colab. Structured for reproducibility and GitHub publication.

A.1 Environment Setup and Data Loading

```
# ——— A.1: Environment Setup ———
# Install required analytical libraries for this project !pip install scikit-learn pandas numpy matplotlib seaborn wordcloud imbalanced-learn

import pandas as pd, numpy as np, matplotlib.pyplot as plt
import seaborn as sns, re, warnings
from wordcloud import WordCloud
warnings.filterwarnings('ignore')

# Mount Google Drive
from google.colab import drive
drive.mount('/content/drive')
BASE = '/content/drive/MyDrive/capstone/'

# Load all four Kaggle datasets
amazon = pd.read_csv(BASE + 'amazon.csv')          # 1,464 records
tweets = pd.read_csv(BASE + 'final_data.csv')      # CrowdFlower 9,093
reviews = pd.read_csv(BASE + 'product_reviews.csv') # 1,000 records
mobile = pd.read_csv(BASE + 'Customer_Tweet_Reviews_of_Mobile_Phone.csv') # 25,000

print(f'Amazon India:      {amazon.shape}')
print(f'CrowdFlower Tweets: {tweets.shape}')
print(f'Product Reviews:   {reviews.shape}')
print(f'Mobile Tweets:     {mobile.shape}')
print(f'Total corpus:      {len(amazon)+len(tweets)+len(reviews)+len(mobile):,} records')
```

The screenshot shows a Google Colab notebook titled 'Untitled0.ipynb'. The left pane displays the code and its output. The output shows that all required libraries are already satisfied, and the datasets are loaded successfully. The right pane shows the file explorer with the 'capstone' folder containing the datasets.

```
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (1.6.1)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (2.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (2.0.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (3.10.0)
Requirement already satisfied: seaborn in /usr/local/lib/python3.12/dist-packages (0.13.2)
Requirement already satisfied: wordcloud in /usr/local/lib/python3.12/dist-packages (1.9.6)
Requirement already satisfied: imbalanced-learn in /usr/local/lib/python3.12/dist-packages (0.14.1)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.12.0)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.5.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2024.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (4.53.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.6)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.1.4)
Requirement already satisfied: sklearn-compat<0.2,>=0.1.5 in /usr/local/lib/python3.12/dist-packages (from imbalanced-learn) (0.1.6)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil) (1.17.0)
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive")
Amazon India:      (1464, 16)
CrowdFlower Tweets: (9093, 3)
Product Reviews:   (1000, 9)
Mobile Tweets:     (25000, 3)
Total corpus:      36,558 records
```

The right pane shows the file explorer with the 'capstone' folder containing the datasets:

- Customer Tweet Reviews of Mobile Phone.csv
- Dataset - Test.csv
- amazon.csv
- final_data.csv
- product_reviews.csv
- s1.csv

A.2 Data Preprocessing

```
# — A.2: Text Preprocessing —
def clean_text(text):
    text = str(text).lower()
    text = re.sub(r'http\S+|www\S+', '', text) # Remove URLs
    text = re.sub(r'<.*?>', '', text) # Remove HTML tags
    text = re.sub(r'[^\w\s]', '', text) # Letters only
    return re.sub(r'\s+', ' ', text).strip()

# Amazon India: derive sentiment from rating
amazon['rating_clean'] = pd.to_numeric(amazon['rating'], errors='coerce')
amazon['sentiment'] = amazon['rating_clean'].apply(
    lambda x: 'Positive' if x >= 4.0 else ('Negative' if x < 3.5 else 'Neutral')
    if pd.notna(x) else None)
amazon['clean'] = amazon['review_content'].apply(clean_text)

# CrowdFlower: map annotated labels to standard classes
def map_emotion(e):
    if 'Positive' in str(e): return 'Positive'
    elif 'Negative' in str(e): return 'Negative'
    elif 'No emotion' in str(e): return 'Neutral'
    else: return None
tweets['label'] = tweets['is there an emotion directed at a brand or product']\
    .apply(map_emotion)
tweets = tweets.dropna(subset=['label'])
tweets['clean'] = tweets['tweet_text'].apply(clean_text)

# Consumer Electronics Reviews
reviews['label'] = reviews['sentiment'].str.capitalize()
reviews['clean'] = reviews['review_text'].apply(clean_text)

# Combine CrowdFlower + Reviews for ML training corpus
combined = pd.concat([tweets[['clean', 'label']], reviews[['clean', 'label']]],
    ignore_index=True).dropna()
combined = combined[combined['label'].isin(['Positive', 'Negative', 'Neutral'])]
print(f'ML Training Corpus: {len(combined):,} records')
print(combined['label'].value_counts())
```

```
def map_emotion(e):
    if 'Positive' in str(e): return 'Positive'
    elif 'Negative' in str(e): return 'Negative'
    elif 'No emotion' in str(e): return 'Neutral'
    else: return None

tweets['label'] = tweets['is there an emotion directed at a brand or product']\
    .apply(map_emotion)
tweets = tweets.dropna(subset=['label'])
tweets['clean'] = tweets['tweet_text'].apply(clean_text)

# Consumer Electronics Reviews
reviews['label'] = reviews['sentiment'].str.capitalize()
reviews['clean'] = reviews['review_text'].apply(clean_text)

# Combine CrowdFlower + Reviews for ML training corpus
combined = pd.concat([tweets[['clean', 'label']], reviews[['clean', 'label']]],
    ignore_index=True).dropna()
combined = combined[combined['label'].isin(['Positive', 'Negative', 'Neutral'])]
print(f'ML Training Corpus: {len(combined):,} records')
print(combined['label'].value_counts())
```

ML Training Corpus: 9,937 records

label	count
Positive	5642
Negative	3562
Neutral	733

Name: count, dtype: int64

```
# — A.3: EDA Visualisations —
fig, axes = plt.subplots(1, 2, figsize=(14, 5))
```

[illegible]

A.4 TF-IDF Feature Engineering and ML Training

```
# — A.4: ML Model Training —
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder, MinMaxScaler

# Encode labels
le = LabelEncoder()
y = le.fit_transform(combined['label'])
X = combined['clean']
print('Label encoding:', dict(zip(le.classes_ , le.transform(le.classes_ ))))

# Stratified train/test split (85% / 15%)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.15, random_state=42, stratify=y)
print(f'Train: {len(X_train):,} | Test: {len(X_test):,}')

# TF-IDF Vectorisation
tfidf = TfidfVectorizer(max_features=5000, ngram_range=(1,2), min_df=2)
X_train_t = tfidf.fit_transform(X_train)
X_test_t = tfidf.transform(X_test)

# — Logistic Regression (Primary Model) —
lr = LogisticRegression(C=1.0, solver='lbfgs', max_iter=500, random_state=42)
lr.fit(X_train_t, y_train)
lr_pred = lr.predict(X_test_t)
print(f'LR Accuracy: {accuracy_score(y_test, lr_pred):.4f}')
print(classification_report(y_test, lr_pred, target_names=le.classes_ ))

# — Random Forest (Comparison) —
rf = RandomForestClassifier(n_estimators=100, max_depth=20, random_state=42)
rf.fit(X_train_t, y_train)
print(f'RF Accuracy: {accuracy_score(y_test, rf.predict(X_test_t)):.4f}')

# — Naive Bayes (Baseline) —
scaler = MinMaxScaler()
nb = MultinomialNB()
nb.fit(scaler.fit_transform(X_train_t.toarray()), y_train)
print(f'NB Accuracy: {accuracy_score(y_test, nb.predict(scaler.transform(X_test_t.toarray()))):.4f}')
```

The screenshot shows a Jupyter Notebook with the following output:

```
Label encoding: {'Negative': np.int64(0), 'Neutral': np.int64(1), 'Positive': np.int64(2)}
Train: 8,446 | Test: 1,491
LR Accuracy: 0.7239
```

	precision	recall	f1-score	support
Negative	0.81	0.19	0.31	110
Neutral	0.72	0.90	0.80	847
Positive	0.74	0.54	0.63	534
accuracy			0.72	1491
macro avg	0.75	0.55	0.58	1491
weighted avg	0.73	0.72	0.70	1491

```
RF Accuracy: 0.6526
NB Accuracy: 0.6989
```

A.5 Evaluation — Confusion Matrix, ROC Curves, Feature Importance



```
# — A.5: Evaluation Plots —
from sklearn.metrics import confusion_matrix, roc_curve, auc
from sklearn.preprocessing import label_binarize
import numpy as np


# Confusion Matrix
cm = confusion_matrix(y_test, lr_pred)
cm_norm = cm.astype(float) / cm.sum(axis=1, keepdims=True)
import seaborn as sns
fig, ax = plt.subplots(figsize=(7, 6))
sns.heatmap(cm_norm, annot=True, fmt='.2f', cmap='Blues',
            xticklabels=le.classes_, yticklabels=le.classes_, ax=ax)
ax.set_title('Normalised Confusion Matrix — Logistic Regression (Data Test Set)',
            fontweight='bold'); ax.set_xlabel('Predicted'); ax.set_ylabel('True')
plt.tight_layout(); plt.savefig('confusion_matrix.png', dpi=150, bbox_inches='tight')
plt.show()

# ROC Curves (One-vs-Rest)
y_test_bin = label_binarize(y_test, classes=list(range(3)))
y_score = lr.predict_proba(X_test_t)
plt.figure(figsize=(9, 6))
for i, (cls, col) in enumerate(zip(le.classes_, ['#E74C3C', '#95A5A6', '#27AE60'])):
    fpr, tpr, _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    plt.plot(fpr, tpr, color=col, lw=2, label=f'{cls} (AUC={auc(fpr,tpr):.3f})')
plt.plot([0,1],[0,1], 'k--', lw=1.5, label='Random Classifier')
plt.xlabel('False Positive Rate'); plt.ylabel('True Positive Rate')
plt.title('ROC Curves — Logistic Regression (One-vs-Rest, Data)', fontweight='bold')
plt.legend(loc='lower right'); plt.grid(alpha=0.3)
plt.savefig('roc_curves.png', dpi=150, bbox_inches='tight'); plt.show()

# Top Feature Importances
feature_names = tfidf.get_feature_names_out()
fig, axes = plt.subplots(1, 2, figsize=(16, 6))
for ax, cls_name, col in zip(axes, ['Positive', 'Negative'], ['#27AE60', '#E74C3C']):
    idx = le.transform([cls_name])[0]
    coefs = lr.coef_[idx]
    top = np.argsort(coefs)[-15:][::-1]
    ax.barh(range(15), coefs[top], color=col, edgecolor='white')
    ax.set_yticks(range(15)); ax.set_yticklabels(feature_names[top], fontsize=9)
    ax.invert_yaxis(); ax.set_title(f'Top 15 {cls_name} Features', fontweight='bold')
    ax.grid(axis='x', alpha=0.3)
plt.tight_layout(); plt.savefig('feature_importance.png', dpi=150, bbox_inches='tight')
plt.show()

# Save Pipeline
import joblib
joblib.dump({'tfidf': tfidf, 'model': lr, 'le': le}, 'sentiment_pipeline.joblib')
print('Pipeline saved to sentiment_pipeline.joblib')
```





A.6 Production Inference Function


```
# — A.6: Production Inference —
def predict_sentiment(text_list, pipeline_path='sentiment_pipeline.joblib'):
    """
    Predicts sentiment for a list of text strings.
    Returns DataFrame with text, cleaned_text, predicted_label, confidence.
    """
    objects = joblib.load(pipeline_path)
    tfidf_v, model, encoder = objects['tfidf'], objects['model'], objects['le']

    cleaned = [clean_text(t) for t in text_list]
    features = tfidf_v.transform(cleaned)
    predictions = model.predict(features)
    probs = model.predict_proba(features)

    return pd.DataFrame({
        'original_text': text_list,
        'predicted_label': encoder.inverse_transform(predictions),
        'confidence': np.max(probs, axis=1).round(3),
        'prob_positive': probs[:, encoder.transform(['Positive'])[0]].round(3),
        'prob_negative': probs[:, encoder.transform(['Negative'])[0]].round(3),
    })

# Test inference on -world style examples
test_texts = [
    'This phone has an absolutely incredible camera — battery life is outstanding!',
    'Terrible product. Stopped working after two weeks. Complete waste of money.',
    'It is okay I guess, nothing special about this device.',
    'boAt earphones are amazing value for the price. Highly recommend!',
    'Mi band stopped tracking after 1 week. App crashes constantly!',
]

results = predict_sentiment(test_texts)
print("\n--- Sample Inference Results ---")
print(results[['original_text', 'predicted_label', 'confidence']].to_string(index=False))
```

The screenshot displays a Jupyter Notebook titled 'Untitled0.ipynb' with a dark theme. The interface includes a top menu bar (File, Edit, View, Insert, Runtime, Tools, Help), a left sidebar with navigation icons, and a right sidebar showing a file explorer. The main code cell contains Python code for sentiment analysis using a pre-trained model. The code defines a function to predict sentiment probabilities and labels for a given text list. It then tests the model on five real-world style examples. The output shows the predicted sentiment labels and confidence scores for each example.

```
[12] ✓ 0s
probs = model.predict_proba(features)

return pd.DataFrame({
    'original_text': text_list,
    'predicted_label': encoder.inverse_transform(predictions),
    'confidence': np.max(probs, axis=1).round(3),
    'prob_positive': probs[:, encoder.transform(['Positive'])[0]].round(3),
    'prob_negative': probs[:, encoder.transform(['Negative'])[0]].round(3),
})

# Test inference on real-world style examples
test_texts = [
    'This phone has an absolutely incredible camera - battery life is outstanding!',
    'Terrible product. Stopped working after two weeks. Complete waste of money.',
    'It is okay I guess, nothing special about this device.',
    'boat earphones are amazing value for the price. Highly recommend!',
    'Mi band stopped tracking after 1 week. App crashes constantly.',
]

results = predict_sentiment(test_texts)
print('\n--- Sample Inference Results ---')
print(results[['original_text', 'predicted_label', 'confidence']].to_string(index=False))

---
-- Sample Inference Results ---
original_text predicted_label confidence
this phone has an absolutely incredible camera - battery life is outstanding! Positive 0.668
Terrible product. Stopped working after two weeks. Complete waste of money. Negative 0.698
It is okay I guess, nothing special about this device. Neutral 0.543
boat earphones are amazing value for the price. Highly recommend! Neutral 0.517
Mi band stopped tracking after 1 week. App crashes constantly. Positive 0.489
```

The file explorer on the right shows a directory structure with files like 'Customer Tweet Reviews of Mobile Phone.csv', 'Dataset - Test.csv', 'amazon.csv', 'final_data.csv', 'product_reviews.csv', and 's1.csv'. The bottom status bar indicates the system time is 11:27 PM and the environment is Python 3.

APPENDIX B: DATASET SAMPLES

B.1 Amazon India Product Dataset — Sample Records

#	Product Name	Rating	Sentiment	Category
1	Wayona Nylon Braided USB to Lightning Cable (3A, 1 Metre)	4.2	Positive	Computers & Accessories Cables
2	Ambrane Unbreakable 60W Fast Charging Cable (TypeC to C)	4.0	Positive	Computers & Accessories Cables
3	pTron Bassbuds Pro In-Ear TWS Earphones with 32Hrs Playtime	3.8	Neutral	Electronics Headphones
4	Samsung Galaxy M04 (Green, 4GB, 64GB Storage) 5000mAh	4.1	Positive	Electronics Smartphones
5	boAt Rockerz 450 Bluetooth Over Ear Headphones	3.9	Neutral	Electronics Headphones
6	Mi Smart Band 6 with 1.56" AMOLED Display	4.0	Positive	Electronics Wearable Technology
7	LG 108 cm (43 Inches) 4K Ultra HD Smart OLED TV	4.3	Positive	Electronics Home Theatre & TV
8	me Buds Air 3S True Wireless Earbuds (Black)	3.5	Neutral	Electronics Headphones
9	Havells Cruise 1000 Watt Mixer Grinder 3 Jars	4.2	Positive	Home & Kitchen Mixer Grinders
10	Redmi 10 Power (8GB RAM 128GB, Sporty Orange)	4.1	Positive	Electronics Smartphones
11	Apple AirPods Pro (2nd Generation) USB-C	4.4	Positive	Electronics Headphones
12	OnePlus Nord CE 3 Lite 5G (128GB, Pastel Lime)	3.7	Neutral	Electronics Smartphones
13	Samsung 43" Crystal iSmart 4K UHD Smart TV 2023	4.2	Positive	Electronics Television
14	Noise ColorFit Pro 4 Hybrid Smart Watch	3.8	Neutral	Electronics Wearable Technology
15	Anker 65W USB C Charger, 4-Port PD Charger Adapter	4.5	Positive	Computers & Accessories Chargers

Table B.1: Representative Amazon India Product Records (Source: Public Dataset — amazon.csv)

B.2 CrowdFlower Tweet Dataset — Sample Records

#	Tweet Text (, Truncated to 80 chars)	Brand	Sentiment
1	Used my iPhone non-stop at the tech conference today — battery drained completely within a few hours. #disappointed	iPhone	Negative
2	Just discovered this amazing iPad/iPhone productivity app — genuinely impressed with the features. Worth checking out!	iPad App	Positive
3	Google Maps keeps crashing on my iPad. Completely unusable right now — need navigation desperately.	Google Maps	Negative
4	Google's latest acquisition is absolutely mind-blowing. Where are they headed next? #innovation #tech	Google	Positive

5	Super excited to launch my new iPad app at the developer meetup tonight! #buildinpublic #ipad #mobiledev	iPad	Positive
6	The updated Google Maps interface is absolutely fantastic — this is exactly what the app needed. Loving every feature!	Google	Positive
7	Photo app keeps freezing when I browse images on my iPad. This Safari browser compatibility issue is ly frustrating.	iPad	Negative
8	My iPhone app just requested access to my photo library without any explanation. Serious privacy concerns here.	iPhone	Negative
9	Got my new iPad today and spent the entire evening gaming on it. Absolutely worth every rupee — display is stunning!	iPad	Positive
10	Put my new iPhone on charge last night and woke up to a full battery. Outstanding battery performance — absolutely no regrets.	iPhone	Positive

Table B.2: Representative CrowdFlower Annotated Tweet Records (Source: Public Dataset — final_data.csv)

B.3 Consumer Electronics Reviews — Sample Records

#	Product	Category	Rating	Sentiment	Feature Mentioned
1	TechPro X20	Smartphones	5	Positive	display
2	TechPro X20	Smartphones	4	Positive	battery life
3	UltraBook Pro 15	Laptops	2	Negative	performance
4	SmartWatch Elite	Wearables	3	Neutral	battery life
5	SoundSphere Speaker	Audio	5	Positive	sound quality
6	HomeAI Assistant	Smart Home	1	Negative	connectivity
7	TechPro X20	Smartphones	4	Positive	camera
8	UltraBook Pro 15	Laptops	5	Positive	build quality
9	SmartWatch Elite	Wearables	2	Negative	accuracy
10	HomeAI Assistant	Smart Home	4	Positive	ease of use

Table B.3: Representative Consumer Electronics Review Records (Source: Public Dataset — product_reviews.csv)

– END OF REPORT –

Plagiarism Report

ORIGINALITY REPORT

%9	%8	%5	%1
SIMILARITY INDEX	INTERNET SOURCES	PUBLICATIONS	STUDENT PAPERS

PRIMARY SOURCES

1	kaggle.com Internet Source	%4
2	arxiv.org Internet Source	%3
3	Hutto, C. J., & Gilbert, E. VADER: A Parsimonious Rule-Based Model for Sentiment Analysis of Social Media Text. Proceedings of ICWSM, 2014 Publication	%1
4	Pang, B., & Lee, L. Opinion Mining and Sentiment Analysis. Foundations and Trends in Information Retrieval, 2008 Publication	%1
5	amazon.in Internet Source	%1
6	Submitted to IIT Delhi Student Paper	%1
7	Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. Proceedings of NAACL, 2019 Publication	%1
8	Pedregosa, F., et al. Scikit-learn: Machine Learning in Python. Journal of Machine Learning Research, 2011 Publication	<%1